

Variational AutoEncoders

LATEST SUBMISSION GRADE

100%

1. For Variational AutoEncoders, which of the following are the correct operations performed in the *latent space*?

1 / 1 point

- ☐ encoder mean * encoder STDev * gaussian distribution
- ☒ encoder mean + encoder STDev * gaussian distribution
- ☐ encoder mean * encoder STDev + gaussian distribution
- ☐ encoder mean + encoder STDev + gaussian distribution

✔ Correct
Correct!

2. Consider the following code, which is used in Variational AutoEncoder to represent the latent space. Fill in the missing piece of code.

1 / 1 point

(Note: Use shape as shape=(batch, dim))


```
class Sampling(tf.keras.layers.Layer):
    def call(self, inputs):
        mu, sigma = inputs
        batch = tf.shape(mu)[0]
        dim = tf.shape(mu)[1]
        epsilon = # YOUR CODE HERE
        return mu + tf.exp(0.5 * sigma) * epsilon
```

```
tf.keras.backend.random_normal(shape=(batch, dim))
```

✔ Correct
Correct!

3. When building the architecture for the decoder for a *convolutional Variational AutoEncoder*, what type of layers will you use? Below is a screenshot of the code with # layer name # written in place of the actual layer that you would use. What goes in place of # layer name #?

1 / 1 point



```
def decoder_layers(inputs, conv_shape):
    units = conv_shape[1] * conv_shape[2] * conv_shape[3]
    x = tf.keras.layers.Dense(units, activation='relu',
                              name="decode_dense1")(inputs)
    x = tf.keras.layers.BatchNormalization()(x)

    x = tf.keras.layers.Reshape((conv_shape[1], conv_shape[2], conv_shape[3]),
                                name="decode_reshape")(x)
    x = tf.keras.layers.# layer name #(filters=64, kernel_size=3, strides=2,
                          padding='same', activation='relu',
                          name="decode_conv2d_2")(x)
    x = tf.keras.layers.BatchNormalization()(x)

    x = tf.keras.layers.# layer name #(filters=32, kernel_size=3, strides=2,
                          padding='same', activation='relu',
                          name="decode_conv2d3")(x)
    x = tf.keras.layers.BatchNormalization()(x)

    x = tf.keras.layers.# layer name #(filters=1, kernel_size=3, strides=1,
                          padding='same', activation='sigmoid', name="decode_final")(x)

    return x
```

- ☐ Conv2D
- ☒ Conv2DTranspose
- ☐ Global AveragePooling2D
- ☐ MaxPooling2D.

✔ Correct
Correct! This will help you invert the convolutional filters applied during encoding.

4. Fill in the missing code for Kullback-Leibler cost function.

1 / 1 point

```
def kl_reconstruction_loss(inputs, outputs, mu, sigma):
    kl_loss = # YOUR CODE HERE
    return tf.reduce_mean(kl_loss) * - 0.5
```

- ☐ $\mu - \text{tf.square}(\sigma) - \text{tf.math.exp}(\mu)$
- ☒ $\text{kl_loss} = 1 + \sigma - \text{tf.square}(\mu) - \text{tf.math.exp}(\sigma)$
- ☐ $\text{kl_loss} = 1 + \mu - \text{tf.square}(\sigma) - \text{tf.math.exp}(\mu)$
- ☐ $\text{kl_loss} = \sigma - \text{tf.square}(\mu) - \text{tf.math.exp}(\sigma)$

✔ Correct
Correct!

5. Which of the following is true regarding *loss* (the variable)?

1 / 1 point

```
for epoch in range(epochs):
    for step, x_batch_train in enumerate(train_dataset):
        with tf.GradientTape() as tape:
            reconstructed = vae(x_batch_train)
            flattened_inputs = tf.reshape(x_batch_train, shape=[-1])
            flattened_outputs = tf.reshape(reconstructed, shape=[-1])
            loss = bce_loss(flattened_inputs, flattened_outputs) * 784
            loss += sum(vae.losses) # Add KLD regularization loss

        grads = tape.gradient(loss, vae.trainable_weights)
        optimizer.apply_gradients(zip(grads, vae.trainable_weights))
```

Correct!

5. Which of the following is true regarding *loss* (the variable)?

1 / 1 point

```
for epoch in range(epochs):
    for step, x_batch_train in enumerate(train_dataset):
        with tf.GradientTape() as tape:
            reconstructed = vae(x_batch_train)
            flattened_inputs = tf.reshape(x_batch_train, shape=[-1])
            flattened_outputs = tf.reshape(reconstructed, shape=[-1])
            loss = bce_loss(flattened_inputs, flattened_outputs) * 784
            loss += sum(vae.losses) # Add KLD regularization loss

        grads = tape.gradient(loss, vae.trainable_weights)
        optimizer.apply_gradients(zip(grads, vae.trainable_weights))
```

- ☐ The closer the values of binary cross entropy loss function are to 0, the closer is the output to expected results.
- ☒ The closer the values of binary cross entropy loss function are to 1, the closer is the output to expected results.

✔ Correct
Correct!